

```
In [22]: import numpy as np
import pandas as pd
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import r2_score, accuracy_score

import warnings
warnings.filterwarnings('ignore')
```

```
In [23]: data = pd.read_csv('diabetes.csv')
data.head()
```

```
Out[23]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
0	6	148	72	35	0	33.6	0.
1	1	85	66	29	0	26.6	0.
2	8	183	64	0	0	23.3	0.
3	1	89	66	23	94	28.1	0.
4	0	137	40	35	168	43.1	2.

```
In [24]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [25]: data.describe()
```

```
Out[25]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diat
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

```
In [26]: data.skew()
```

```
Out[26]: Pregnancies          0.901674
Glucose          0.173754
BloodPressure    -1.843608
SkinThickness    0.109372
Insulin          2.272251
BMI              -0.428982
DiabetesPedigreeFunction  1.919911
Age              1.129597
Outcome          0.635017
dtype: float64
```

```
In [27]: data.kurt()
```

```
Out[27]: Pregnancies          0.159220
Glucose          0.640780
BloodPressure     5.180157
SkinThickness    -0.520072
Insulin          7.214260
BMI              3.290443
DiabetesPedigreeFunction  5.594954
Age              0.643159
Outcome          -1.600930
dtype: float64
```

```
In [28]: data.mode().iloc[0]
```

```
Out[28]: Pregnancies          1.000
Glucose          99.000
BloodPressure     70.000
SkinThickness      0.000
Insulin           0.000
BMI              32.000
DiabetesPedigreeFunction  0.254
Age              22.000
Outcome           0.000
Name: 0, dtype: float64
```

```
In [29]: X = data.drop('Outcome', axis=1)
y = data['Outcome']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
```

```
In [30]: #Linear Regression
lin_re = LinearRegression()
lin_re.fit(X_train, y_train)
y_pred_lin = lin_re.predict(X_test)
r = r2_score(y_test, y_pred_lin)

print('R-squared is:', r)
```

R-squared is: 0.25500281176741757

```
In [31]: #Logistic Regression
log_re = LogisticRegression()
log_re.fit(X_train, y_train)
y_pred_log = log_re.predict(X_test)

a = accuracy_score(y_test, y_pred_log)
print('Accuracy is:', a)
```

Accuracy is: 0.7532467532467533

```
In [ ]:
```