

Practical 1

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.metrics import accuracy_score , confusion_matrix
```

```
In [2]: data = pd.read_csv ("iris.csv")
```

```
In [3]: data.head()
```

```
Out[3]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
In [4]: x = data.iloc[:, :-1]
y = data['Species']
```

```
In [5]: x_test, x_train , y_test, y_train = train_test_split (x, y , test_size = 0.
```

```
In [6]: scale = StandardScaler()
```

```
In [8]: x_train_s = scale.fit_transform(x_train)
x_test_s = scale.transform(x_test)
```

```
In [9]: lda = LinearDiscriminantAnalysis()
```

```
In [12]: x_train_l = lda.fit_transform(x_train_s , y_train)
x_test_l = lda.transform(x_test_s)
```

```
In [13]: classify = LogisticRegression()
```

```
In [14]: classify.fit(x_train_l , y_train)
y_pred = classify.predict(x_test_l)
```

```
In [16]: accuracy = accuracy_score(y_test, y_pred)
conf_m = confusion_matrix(y_test, y_pred)

print("Accuracy :", accuracy)
print ("Confusion Matrix : \n", conf_m)
```

```
Accuracy : 1.0
Confusion Matrix :
[[40  0  0]
 [ 0 41  0]
 [ 0  0 39]]
```

Practical 2

```
In [1]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

```
In [3]: import warnings
warnings.filterwarnings("ignore")
```

```
In [5]: data = pd.read_csv("diabetes.csv")
data.head()
```

```
Out[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
0	6	148	72	35	0	33.6	0.
1	1	85	66	29	0	26.6	0.
2	8	183	64	0	0	23.3	0.
3	1	89	66	23	94	28.1	0.
4	0	137	40	35	168	43.1	2.

```
In [6]: data.skew()
```

```
Out[6]: Pregnancies      0.901674
Glucose      0.173754
BloodPressure -1.843608
SkinThickness 0.109372
Insulin      2.272251
BMI          -0.428982
DiabetesPedigreeFunction 1.919911
Age          1.129597
Outcome      0.635017
dtype: float64
```

```
In [7]: data.kurt()
```

```
Out[7]: Pregnancies      0.159220  
Glucose      0.640780  
BloodPressure  5.180157  
SkinThickness -0.520072  
Insulin      7.214260  
BMI          3.290443  
DiabetesPedigreeFunction  5.594954  
Age          0.643159  
Outcome      -1.600930  
dtype: float64
```

```
In [8]: data.mode().iloc[0]
```

```
Out[8]: Pregnancies      1.000  
Glucose      99.000  
BloodPressure  70.000  
SkinThickness  0.000  
Insulin      0.000  
BMI          32.000  
DiabetesPedigreeFunction  0.254  
Age          22.000  
Outcome      0.000  
Name: 0, dtype: float64
```

```
In [14]: x = data.drop('Outcome', axis=1)  
y = data['Outcome']
```

```
In [15]: x_test, x_train , y_test, y_train = train_test_split(x, y, test_size = 0.2,
```

```
In [16]: linear_r = LinearRegression()  
linear_r.fit(x_train, y_train)  
y_pred_lin = linear_r.predict(x_test)
```

```
In [17]: r2_score = r2_score(y_test, y_pred_lin)  
print(r2_score)
```

```
0.15432083745026837
```

```
In [18]: log_r = LogisticRegression()  
log_r.fit(x_train , y_train)  
y_pred_log = log_r.predict(x_test)
```

```
In [19]: accuracy = accuracy_score(y_test, y_pred_log)  
print("Accuracy :", accuracy)
```

```
Accuracy : 0.749185667752443
```

Practical 3

```
In [55]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import precision_score , recall_score
```

```
In [26]: data3 = pd.read_csv("Social_Network_Ads.csv")
data3.head()
```

```
Out[26]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [28]: x = data3.iloc[:, [2,3]]
y = data3["Purchased"]
```

```
In [41]: x_train , x_test, y_train, y_test = train_test_split(x,y, test_size = 0.2 ,
```

```
In [45]: std_scale = StandardScaler()
x_train_s = std_scale.fit_transform(x_train)
x_test_s = std_scale.transform(x_test)
```

```
In [46]: k = 5
knn = KNeighborsClassifier(n_neighbors = k)
knn.fit(x_train_s , y_train)
```

```
Out[46]: KNeighborsClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [47]: y_pred = knn.predict(x_test_s)
```

```
In [49]: conf_m = confusion_matrix(y_test, y_pred)
print("Conf matrix: \n", conf_m)
```

```
Conf matrix:
[[48  4]
 [ 3 25]]
```

```
In [50]: accuracy = accuracy_score(y_test, y_pred)
print("Accuracy :", accuracy)
```

Accuracy : 0.9125

```
In [51]: error_rate = 1 - accuracy
print(error_rate)
```

0.08750000000000002

```
In [56]: precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
```

```
In [57]: print("Precision score :", precision)
print("recall score :", recall)
```

Precision score : 0.8620689655172413
recall score : 0.8928571428571429

Practical 4

```
In [58]: import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
```

```
In [59]: data4 = pd.read_csv("iris.csv")
data4.head()
```


```
Out[59]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
In [60]: x = data.iloc[:, [1,2,3,4]]
```

```
In [65]: inertia = []
for i in range(1,11):
    kmeans = KMeans(n_clusters = i , max_iter = 300 , random_state = 42)
    kmeans.fit(x)
    inertia.append(kmeans.inertia_)
```

```
In [66]: plt.plot(range(1,11), inertia, marker = 'o')
plt.xlabel("No of clusters")
plt.ylabel("Inertia")
plt.title("Elbow Method")
plt.show()
```



Practical 5

```
In [16]: from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder
```

```
In [17]: data5 = pd.read_csv("car_evaluation.csv")
data5.head()
```

```
Out[17]:
```

	Buying price	Maintance cost	No of doors	No of persons	lug_boot	safety	Decision
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
In [18]: col_names = ['Buying price', 'Maintance cost', 'No of doors', 'No of person']
data5.columns = col_names

data5.head()
```

```
Out[18]:
```

	Buying price	Maintance cost	No of doors	No of persons	lug_boot	safety	Decision
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
In [20]: le = LabelEncoder()

for col in data5.columns:
    if data5[col].dtype != 'object':
        continue

    data5[col] = le.fit_transform(data5[col])

data5.head()
```

```
Out[20]:
```

	Buying price	Maintance cost	No of doors	No of persons	lug_boot	safety	Decision
0	3	3	0	0	2	1	2
1	3	3	0	0	2	2	2
2	3	3	0	0	2	0	2
3	3	3	0	0	1	1	2
4	3	3	0	0	1	2	2

```
In [23]: x = data5.drop(['Decision'], axis = 1)
y = data5['Decision']

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3)
x_train.shape, x_test.shape
```

```
Out[23]: ((1209, 6), (519, 6))
```

```
In [24]: rfc = RandomForestClassifier(n_estimators = 18+25)
rfc.fit(x_train,y_train)
```

```
Out[24]: RandomForestClassifier(n_estimators=43)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [26]: y_pred = rfc.predict(x_test)
accuracy = accuracy_score(y_test,y_pred)
conf_matrix = confusion_matrix(y_test,y_pred)

print('Accuracy:',accuracy*100,'%','\n')
print('Confusion Matrix:\n',conf_matrix)
```

Accuracy: 96.91714836223507 %

Confusion Matrix:

[[108	7	2	1]
[0	18	0	1]
[2	0	356	0]
[3	0	0	21]]

Practical 6

```
In [15]: maze = np.array([
    [0,0,0,0,0],
    [0,1,0,1,0],
    [0,0,0,0,0],
    [0,1,1,1,0],
    [0,0,0,0,2]
])

learning_rate = 0.1
discount_factor = 0.9
epsilon = 0.1
num_episodes = 1000

num_states, num_actions = maze.size, 4
Q = np.zeros((num_states, num_actions))

for _ in range(num_episodes):
    state = 0

    while True:
        action = np.random.choice(num_actions) if np.random.uniform(0,1) <
        new_state = state + [0,1,2,3][action]
        reward = [-1,1,0][maze.flat[new_state]]
        if reward: break
        state = new_state

    current_state = 0
    while current_state != 16:
        action = np.argmax(Q[current_state, :])
        current_state = current_state + (action + 1)
        print("Agent moved to state :", current_state)
```

```
Agent moved to state : 1
Agent moved to state : 2
Agent moved to state : 3
Agent moved to state : 4
Agent moved to state : 5
Agent moved to state : 6
Agent moved to state : 7
Agent moved to state : 8
Agent moved to state : 9
Agent moved to state : 10
Agent moved to state : 11
Agent moved to state : 12
Agent moved to state : 13
Agent moved to state : 14
Agent moved to state : 15
Agent moved to state : 16
```

In []:

